

Two release risks. One targeted path to measure and fix them.

OPENAI / OBSERVED RELEASE REGRESSION

-16%

Sol 26/44 vs GPT-5.5 75.0%.

STRICT COMPLETION REGRESSION

-16%

Sol 59.1% vs GPT-5.5 75.0%

26/44 VS 33/44

NET FAILURES

+7

Sol 18 vs GPT-5.5 11

7 FEWER COMPLETIONS

BASIS DIAGNOSIS

14 vs 7

31.8% vs 15.9%

2x RATE

- Sol completed 26 of 44 valid finance questions strictly, versus 33 of 44 for GPT-5.5.
- That seven-answer regression is the release signal. It is a narrow finance evaluation, not a general model ranking.
- Basis selection is the diagnosis, not the headline: 14 misses versus 7, while completeness and no-submission counts did not move.

INSPECT OPENAI TRACE

ANTHROPIC / DIFFERENTIATED FAILURE SIGNAL

Fable trails GPT-5.5 on the overall frontier benchmark.

Fable 70.5% versus GPT-5.5 75.0% across the same valid 44.

STRICT COMPLETION GAP

-4.5%

Fable 70.5% vs GPT-5.5 75.0%

31/44 VS 33/44

TOTAL FAILURES

13

Fable 13/44 vs GPT-5.5 11/44

+2 ANSWERS

COMPLETENESS

5

Fable 5 + other 8 | GPT-5.5 3 + other 8

+66.7% = 2-FAILURE GAP

- Fable trails GPT-5.5 by two strict answers across the 44 valid finance questions.
- Those two additional misses are both completeness. All other failure types total eight for each model.
- These are polished answers that omit a decision-critical field. The target is required-field recall, not length.

INSPECT ANTHROPIC TRACE

WHY TURING

Turn the observed failure into a trainable, auditable program.

01 MODEL TRACE

Locate the exact decision failure.

02 FINANCE EXPERT SPEC

Freeze basis, fields, evidence, and contradictions.

03 ATOMIC VERIFIER

Create a training-ready signal without rewarding verbosity.

04 SEALED RESULT

Measure lift, regressions, exploits, and transfer.

Turing supplies expert tasks, keys, verifiers, adjudication, and the sealed evaluation. The lab owns the model update.

FIRST STEP: SCOPE ONE FAILURE FAMILY AND FREEZE THE SUCCESS GATE

Valid 44 removes six confirmed benchmark key or specification defects from the 50 attempted tasks. Strict success uses all 44 valid attempts, so no-submissions count against success and remain separately labeled: one each for GPT-5.5 and Sol. Single-run release signal, not a general ranking.