

Two release risks. One targeted path to measure and fix them.

OPENAI / OBSERVED RELEASE REGRESSION

GPT-5.6 Sol's basis-selection failure rate was 31.8%.

GPT-5.5 baseline: 15.9% across the same 44 valid questions.

BASIS FAILURE RATE

31.8%

Sol 14/44 vs GPT-5.5 7/44

2x RATE

STRICT SUCCESS

59.1%

Sol 26/44 vs GPT-5.5 33/44

-15.9 POINTS

NET FAILURES

+7

Basis 14/7 | Complete 3/3 |
No-submit 1/1

18 VS 11

- The newer checkpoint delivered seven fewer strict successes across the 44 valid finance questions.
- The entire seven-failure increase is basis selection. Completeness and no-submission counts did not move.
- That makes this a targetable release risk: train and retest basis discipline without rewarding blind copying.

INSPECT OPENAI TRACE

ANTHROPIC / DIFFERENTIATED FAILURE SIGNAL

Fable trails GPT-5.5 on the overall frontier benchmark.

Fable 70.5% versus GPT-5.5 75.0% across the same valid 44.

STRICT SUCCESS

70.5%

Fable 31/44 vs GPT-5.5 33/44

-4.5 POINTS

TOTAL FAILURES

13

Fable 13/44 vs GPT-5.5 11/44

+2 ANSWERS

COMPLETENESS

5

Fable 5 + other 8 | GPT-5.5 3 +
other 8

+66.7% = 2-FAILURE GAP

- Fable trails GPT-5.5 by two strict answers across the 44 valid finance questions.
- Those two additional misses are both completeness. All other failure types total eight for each model.
- These are polished answers that omit a decision-critical field. The target is required-field recall, not length.

INSPECT ANTHROPIC TRACE

WHY TURING

Turn the observed failure into a trainable, auditable program.

01 MODEL TRACE

Locate the exact decision failure.

02 FINANCE EXPERT SPEC

Freeze basis, fields, evidence, and contradictions.

03 ATOMIC VERIFIER

Create a training-ready signal without rewarding verbosity.

04 SEALED RESULT

Measure lift, regressions, exploits, and transfer.

Turing supplies expert tasks, keys, verifiers, adjudication, and the sealed evaluation. The lab owns the model update.

FIRST STEP: SCOPE ONE FAILURE FAMILY AND FREEZE THE SUCCESS GATE

Valid 44 removes six confirmed benchmark key or specification defects from the 50 attempted tasks. Strict success uses all 44 valid attempts, so no-submissions count against success and remain separately labeled: one each for GPT-5.5 and Sol. Single-run release signal, not a general ranking.