Two release risks. One targeted path to measure and fix them.

OPENAI / OBSERVED RELEASE REGRESSION

GPT-5.6 Sol doubled observed basis-selection failures.

7 TO 14

2x COUNT

Basis-selection failures / 49 graded answers each

Matched strict: 32/48 to 26/48

Matched basis: 7 to 13

- GPT-5.6 Sol produced twice as many basis-selection failures as GPT-5.5 on each model's full 49-grade run: 14 versus 7.
- On the identical 48 graded tasks, strict completion moved from 32 of 48 to 26 of 48.
- The added misses concentrate in specification adherence after retrieval: wrong metric, period, filing vintage, or convention.

INSPECT OPENAI TRACE

ANTHROPIC / DIFFERENTIATED FAILURE SIGNAL

Fable's standout risk is an incomplete answer that looks finished.

COMPLETENESS FAILURES / IDENTICAL 48 GRADED TASKS

- Fable produced five completeness failures, versus three for GPT-5.5 and three for GPT-5.6 Sol.
- These answers often found most of the facts, then omitted one decision-critical comparison or term.
- The target is required-field recall, not longer answers. Unsupported claims and contradictions stay as guardrails.

INSPECT ANTHROPIC TRACE

WHY TURING

Turn the observed failure into a trainable, auditable program.

01 MODEL TRACE
Locate the exact decision failure.

02 FINANCE EXPERT SPEC
Freeze basis, fields, evidence, and contradictions.

03 ATOMIC VERIFIER
Create a training-ready signal without rewarding verbosity.

04 SEALED RESULT
Measure lift, regressions, exploits, and transfer.

Turing supplies expert tasks, keys, verifiers, adjudication, and the sealed evaluation. The lab owns the model update.

FIRST STEP: SCOPE ONE FAILURE FAMILY AND FREEZE THE SUCCESS GATE

Single-run evaluation; two no-submission tasks were retried. Graded bases: 49, 50, 49. Direct comparisons use the 48 tasks graded by all three. Six shared benchmark defects are separated from model behavior. Sol versus GPT-5.5 touches zero at the paired interval boundary.