


Both frontiers converge at 60.0% on agentic finance work.

GPT-5.5: 12/20. Claude Fable 5: 12/20. Identical 20 tasks, my grader, my runs. Task-level agreement 18/20: the frontier fails the SAME work.



- 8 models, 174 graded completions, one validated judge; ladder spans 12% to 63% strict on graded completions (6% at the floor over everything attempted).
- Best public leaderboard number: under 58% `Vals v2, external`. Human expert baseline: \$25.66 and 16.8 min per query `external`.

The agent had the answer at step 15. It never said it.

Gemini 3.5 Flash, task 7: "What was the total consideration TKO paid for the Endeavor assets?"

- 1 Steps 1-14: competent search; finds the right 8-K.
- 2 Step 15: retrieves the consideration verbatim from the 8-K: **"\$3.30 billion (measured using the volume-weighted...)"** (the key's \$3.25B closing consideration plus the \$50M adjustment).
- 3 Steps 16-49: thirty-four more turns re-searching its own answer; the query strings literally contain *"3.30 billion" OR "3.3 billion" OR "3,300 million"*.
- 4 Step 50: turn cap. No submission. Graded fail. 317 seconds, 50 tool calls, the answer in hand about 110 seconds in.

The gap is not finance knowledge. It is knowing when you are done: verify-then-commit is a trainable behavior with a deterministic reward. The same give-up family dominates Gemini 2.5 Flash's 25-of-50 no-answer run, and a Gemini 2.5 Pro probe on that model's failed tasks completed 14/15 on the same harness: loop discipline separates even sibling models.

Every failure mode prices a product.

WHAT I MEASURED (MY INSTANCE COUNTS)	WHAT TURING BUILDS
retrieval-miss (19) + step-limit-giveup (9)	RL environments over live EDGAR/web tooling with verifiable rewards: the reward is submit-when-verified, escalate-when-stuck. My 2.5 Pro-vs-2.5 Flash probe shows this is trainable competence, not scale.
methodology-mismatch (12) + arithmetic-or-units (3)	Expert-authored rubric packs for financial-modeling conventions (windows, base years, average-vs-ending, seasonality). Deterministic pass/fail checks per convention: the RFT-shaped reward.
partial-answer (17) + format-noncompliance (2)	Completeness-graded eval sets: conjunction rubrics that price every sub-claim, so silent omission stops being free.
hallucinated-negative-disclosure (1) + hallucinated-tool-completion (1) + context-loss (1)	Attribution-graded training data: every claim priced against the agent's own retrieved evidence. Low counts, catastrophic in production; exactly what banks will ask about.
stale-reference-answer (5 models hit) + mislabeled-input (1, smoke-corroborated) + divergent public copies (12 of 49 records)	Quarterly refresh + expert adjudication pipeline for eval corpora. My two data-QA findings are the live demo that eval assets decay and disagree without it.

The dataset row is measured, not asserted: a stale key graded down five correct models, and the two public copies agree verbatim on only 42 of 50 questions (12 of 49 pairable records differ in key or rubric). Refresh + expert adjudication is a product line, demonstrated live on a respected benchmark.

The same expert work is the eval AND the training signal.

- Every rubric check here is deterministic, atomic, pass/fail: graded conjunctions over expert-written criteria.
- That is exactly the shape an RFT reward needs: no reward model to game, and lucky guesses do not survive multi-criterion conjunction.
- First two reward targets from my taxonomy: submit-when-verified (the give-up modes) and convention compliance (windows, base years, average-vs-ending).

I didn't run RL. The grader is the right SHAPE for an RFT reward: deterministic, atomic, pass/fail. Proving convergence is the pilot Turing would run with you.

Frontier inference averaged \$0.31-1.01 per completed task in my runs (per-run means) against a \$25.66 human baseline `external`; the capability gap is the last thing standing between those two numbers.

Vals proved the market. Turing is built for what labs need next.

- **Credit where due:** the benchmark discriminates, and Vals' v2 private set `external` proves labs pay for private finance evals.
- **The governance gap:** one vendor running the public leaderboard AND licensing the private eval is structurally conflicted. Labs need an eval/train wall from a neutral party with no leaderboard to defend.
- **What my evidence says to buy:** exclusive slices sized to YOUR weak categories (the convergence finding names them), quarterly refresh (the stale-key finding), expert adjudication (the copy-divergence finding: 12 of 49 records), RL environments on live EDGAR tooling (the 2.5 Pro-vs-2.5 Flash probe).

The ask: a scoping pilot. Two categories where your internal evals hurt most; Turing delivers a sized, adjudicated private slice plus an RFT-reward pilot on the give-up modes. My grader, taxonomy, and findings are the starting spec.