SLIDE 1 / THESIS

Frontier models can find the right filing and still submit the wrong financial answer.

The same 20 live tasks put GPT-5.6 Sol at 9/20 strict passes and GPT-5.5 at 12/20. Sol agrees with each incumbent frontier on 17/20 tasks. This is a bounded, task-level observation from one paired slice, not a claim that the models are equivalent or stable.

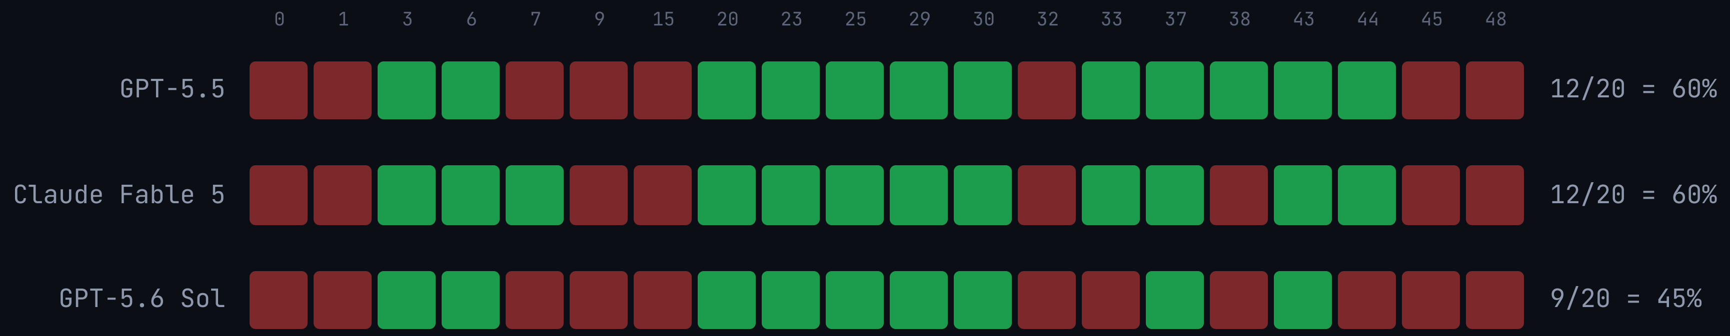

- 9 models, 194 graded completions, one validated judge; ladder spans 12% to 63% strict on graded completions (6% at the floor over everything attempted).
- Best public leaderboard number: under 58% `Vals v2, external`. Human expert baseline: $25.66 and 16.8 min per query `external`.

The agent had the answer at step 15. It never said it.

Gemini 3.5 Flash, task 7: "What was the total consideration TKO paid for the Endeavor assets?"

1. Steps 1-14: competent search; finds the right 8-K.
2. Step 15: retrieves the consideration verbatim from the 8-K: **"$3.30 billion (measured using the volume-weighted...)"** (the key's $3.25B closing consideration plus the $50M adjustment).
3. Steps 16-49: thirty-four more turns re-searching its own answer; the query strings literally contain *"3.30 billion" OR "3.3 billion" OR "3,300 million"*.
4. Step 50: turn cap. No submission. Graded fail. 317 seconds, 50 tool calls, the answer in hand about 110 seconds in (wall clock).

The trace shows a submission failure after relevant evidence entered context. It motivates an outcome-verifier hypothesis. The same give-up family appears in Gemini 2.5 Flash's 25-of-50 no-answer run, while a Gemini 2.5 Pro probe completed 14/15 on the same harness. A sealed pilot would test whether a reward changes that behavior.

Every failure classified.

WHAT I MEASURED (MY INSTANCE COUNTS)	WHAT TURING BUILDS
Evidence retrieval (19) + Turn-limit give-up (9)	RL environments over live EDGAR/web tooling with verifiable rewards: the reward is submit-when-verified, escalate-when-stuck. My 2.5 Pro-vs-2.5 Flash probe shows this is trainable competence, not scale.
Basis selection (12) + Arithmetic or units (3)	Expert-authored rubric packs for financial-modeling conventions (windows, base years, average-vs-ending, seasonality). Deterministic pass/fail checks per convention: the RFT-shaped reward.
Completeness (17) + Output format (2)	Completeness-graded eval sets: conjunction rubrics that price every sub-claim, so silent omission stops being free.
Negative disclosure hallucination (1) + False tool-completion belief (1) + Context loss (1)	Attribution-graded training data: every claim priced against the agent's own retrieved evidence. Low counts, catastrophic in production; exactly what banks will ask about.
Key freshness (5 models hit) + Input-label defect (1, smoke-corroborated) + divergent public copies (12 of 49 records)	Quarterly refresh + expert adjudication pipeline for eval corpora. My two data-QA findings are the live demo that eval assets decay and disagree without it.

The dataset row is measured, not asserted: the audit found a stale key and differences between two public copies. That is a bounded QA finding that motivates versioned keys, disclosed adjudication, and a sensitivity check before model decisions rely on small deltas.

The same expert work is the eval AND the training signal.

- Every rubric check here is deterministic, atomic, pass/fail: graded conjunctions over expert-written criteria.
- That shape can become an RFT reward, but it still needs adversarial testing for overfitting, reward hacking, and regressions on legitimate recomputation.
- First candidate reward targets from the taxonomy: submit-when-verified and convention compliance across windows, base years, and average-versus-ending.

> *The grader is shaped like a training signal: deterministic, atomic, pass/fail. The pilot tests improvement on a sealed slice.*

Frontier inference averaged \$0.31-\$1.01 per attempted task in my runs, versus a \$25.66 published benchmark-execution baseline `external`. This is not a deployed total-cost-of-ownership or labor-substitution claim.

Independent evaluation governance for finance-agent releases.

- **What this audit shows:** a finance benchmark needs versioned keys, disclosed adjudication, and a separation between model development and evaluation maintenance.
- **What Turing supplies:** private slices sized to observed failure families, expert rubric authorship, key refresh, and a sealed eval wall.
- **What the pilot decides:** whether an independently authored verifier improves the selected behavior on fresh finance tasks without breaking adjacent work.

The ask: a scoped diagnostic. Choose two failure families, agree the pass gate, and build the independently adjudicated baseline and sealed evaluation plan. Training is a gated phase once access, prevalence, and generalization criteria are met.